

```

import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras import datasets, layers, models
from tensorflow.keras.utils import to_categorical
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns

# Set random seed for reproducibility
np.random.seed(42)
tf.random.set_seed(42)

# a. Perform Data Pre-processing
# Load and preprocess the MNIST dataset
(train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()

# Normalize pixel values to be between 0 and 1
train_images = train_images.astype('float32') / 255.0
test_images = test_images.astype('float32') / 255.0

# Reshape images to include channel dimension (MNIST images are grayscale)
train_images = train_images.reshape(train_images.shape[0], 28, 28, 1)
test_images = test_images.reshape(test_images.shape[0], 28, 28, 1)

# One-hot encode the labels
train_labels = to_categorical(train_labels, 10)
test_labels_one_hot = to_categorical(test_labels, 10)

# Display sample images
plt.figure(figsize=(10, 5))
for i in range(10):
    plt.subplot(2, 5, i+1)
    plt.imshow(train_images[i].reshape(28, 28), cmap='gray')
    plt.title(f"Label: {np.argmax(train_labels[i])}")
    plt.axis('off')
plt.tight_layout()
plt.savefig('mnist_samples.png')
plt.close()

# b. Define Model and perform training
def create_model():
    model = models.Sequential([
        # First Convolutional Layer
        layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
        layers.MaxPooling2D((2, 2)),

        # Second Convolutional Layer
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.MaxPooling2D((2, 2)),

        # Third Convolutional Layer
        layers.Conv2D(64, (3, 3), activation='relu'),

        # Flatten and Dense layers
        layers.Flatten(),
        layers.Dense(128, activation='relu'),
        layers.Dropout(0.5), # Dropout for regularization
        layers.Dense(10, activation='softmax') # 10 output classes for digits 0-9
    ])

    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    return model

# Create and train the model
model = create_model()
model.summary()

# Train the model with validation split
history = model.fit(train_images, train_labels, epochs=10,
                    batch_size=128, validation_split=0.2,
                    verbose=1)

# Plot training history
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')

```

```

plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.savefig('training_history.png')
plt.close()

# c. Evaluate Results using confusion matrix
# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels_one_hot)
print(f'Test accuracy: {test_acc:.4f}')

# Make predictions
predictions = model.predict(test_images)
predicted_classes = np.argmax(predictions, axis=1)
true_classes = np.argmax(test_labels_one_hot, axis=1)

# Create confusion matrix
cm = confusion_matrix(true_classes, predicted_classes)
print("Confusion Matrix:")
print(cm)

# Plot confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=list(range(10)),
            yticklabels=list(range(10)))
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.savefig('confusion_matrix.png')
plt.close()

# Print classification report
print("\nClassification Report:")
print(classification_report(true_classes, predicted_classes))

# Visualize some misclassified examples
misclassified_idx = np.where(predicted_classes != true_classes)[0]
if len(misclassified_idx) > 0:
    plt.figure(figsize=(12, 4))
    for i, idx in enumerate(misclassified_idx[:10]):
        if i < 10: # Display up to 10 misclassified examples
            plt.subplot(2, 5, i+1)
            plt.imshow(test_images[idx].reshape(28, 28), cmap='gray')
            plt.title(f"True: {true_classes[idx]}, Pred: {predicted_classes[idx]}")
            plt.axis('off')
    plt.tight_layout()
    plt.savefig('misclassified_examples.png')
    plt.close()

print("Analysis completed.")

```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
 11490434/11490434 — 0s 0us/step
 /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` argument to `Conv2D` layers.
 super().__init__(activity_regularizer=activity_regularizer, **kwargs)
 Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36,928
flatten (Flatten)	(None, 576)	0
dense (Dense)	(None, 128)	73,856
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1,290

Total params: 130,890 (511.29 KB)

Trainable params: 130,890 (511.29 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

375/375 — 43s 109ms/step - accuracy: 0.7614 - loss: 0.7353 - val_accuracy: 0.9766 - val_loss: 0.0753

Epoch 2/10

375/375 — 81s 108ms/step - accuracy: 0.9681 - loss: 0.1077 - val_accuracy: 0.9833 - val_loss: 0.0547

Epoch 3/10

375/375 — 39s 102ms/step - accuracy: 0.9792 - loss: 0.0694 - val_accuracy: 0.9878 - val_loss: 0.0423

Epoch 4/10

375/375 — 41s 103ms/step - accuracy: 0.9843 - loss: 0.0552 - val_accuracy: 0.9887 - val_loss: 0.0414

Epoch 5/10

375/375 — 41s 102ms/step - accuracy: 0.9858 - loss: 0.0442 - val_accuracy: 0.9902 - val_loss: 0.0407

Epoch 6/10

375/375 — 42s 105ms/step - accuracy: 0.9883 - loss: 0.0387 - val_accuracy: 0.9890 - val_loss: 0.0419

Epoch 7/10

375/375 — 45s 117ms/step - accuracy: 0.9892 - loss: 0.0338 - val_accuracy: 0.9902 - val_loss: 0.0374

Epoch 8/10

375/375 — 77s 105ms/step - accuracy: 0.9913 - loss: 0.0284 - val_accuracy: 0.9909 - val_loss: 0.0363

Epoch 9/10

375/375 — 41s 106ms/step - accuracy: 0.9920 - loss: 0.0266 - val_accuracy: 0.9911 - val_loss: 0.0358

Epoch 10/10

375/375 — 41s 107ms/step - accuracy: 0.9940 - loss: 0.0186 - val_accuracy: 0.9918 - val_loss: 0.0359

313/313 — 4s 12ms/step - accuracy: 0.9886 - loss: 0.0380

Test accuracy: 0.9917

313/313 — 3s 8ms/step

Confusion Matrix:

```
[
  [ 976  1  0  0  0  0  1  1  1  0],
  [  0 1132  0  1  0  2  0  0  0  0],
  [  2  2 1024  0  0  0  0  3  1  0],
  [  0  0  1 1005  0  3  0  1  0  0],
  [  0  0  0  0  973  0  3  0  0  6],
  [  1  0  0  3  0 886  1  1  0  0],
  [  2  2  0  0  1  5 947  0  1  0],
  [  0  4  3  4  0  0  0 1016  0  1],
  [  1  0  1  2  0  0  0  3 965  2],
  [  4  1  0  1  2  2  0  4  2 993]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.99	1.00	0.99	980
1	0.99	1.00	0.99	1135
2	1.00	0.99	0.99	1032
3	0.99	1.00	0.99	1010
4	1.00	0.99	0.99	982
5	0.99	0.99	0.99	892
6	0.99	0.99	0.99	958
7	0.99	0.99	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.98	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000